Operationalizing an AWS ML Project

1. Initial Setup

   I have chosen the "ml.t3.medium" instance type for this Notebook due to a variety of considerations, all of which cater to the specific needs of this project.

   Firstly, the nature of our project doesn't necessitate the employment of a high-performance CPU or an abundance of RAM for the successful execution of our Jupyter notebooks. Rather than focusing on raw computational power, it's more prudent to consider the time duration our notebook instance will remain active.

   Considering the potentially substantial duration of our project, it's essential to select an instance with a cost-effective hourly rate, in addition to providing a reasonable level of CPU and RAM. This strategic decision is crucial to curbing excessive costs while ensuring the functionality and efficiency of our project are uncompromised.

   In light of these considerations, the "ml.t3.medium" instance emerges as an optimal choice. It provides a balance between cost-efficiency and computational prowess that fits our project's requirements. Despite a slower startup time, this instance type is more cost-effective per hour, a trade-off that is acceptable given our project's lack of dependency on instant start-up.

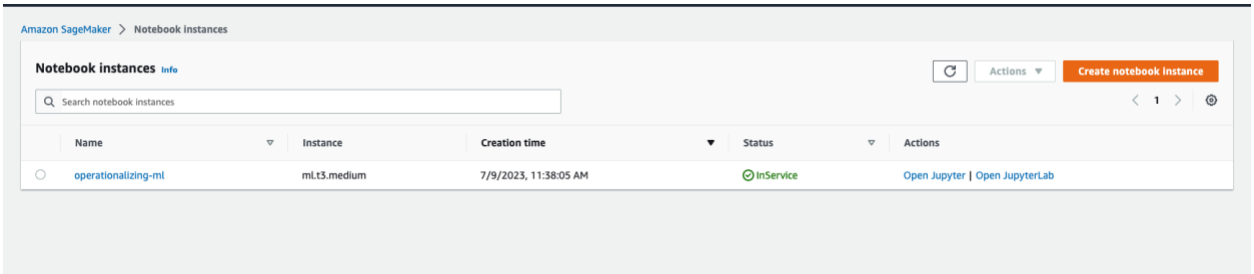

S3 Bucket

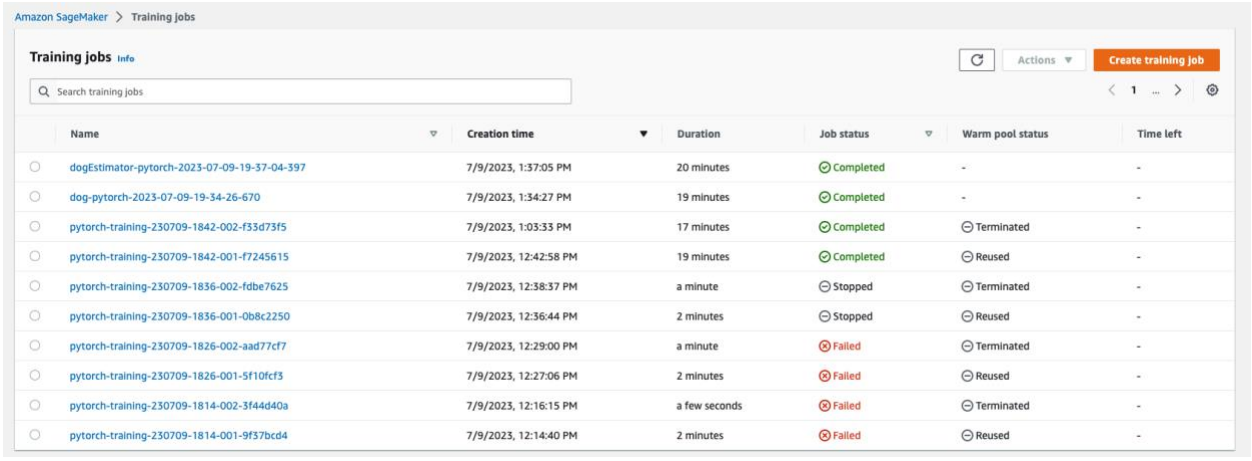

Training and Tuning Jobs

Training and Tuning didn't take that long because of the instance type I used.

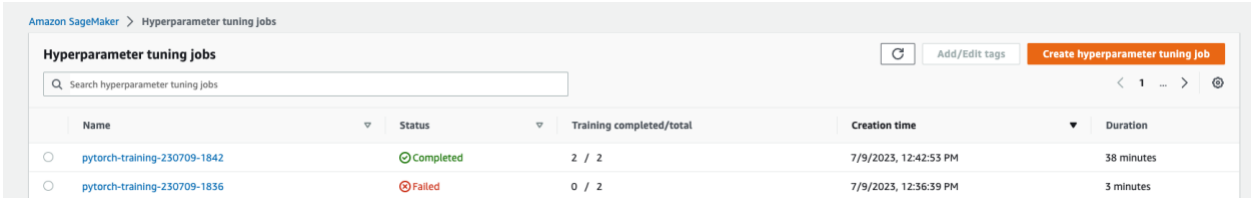

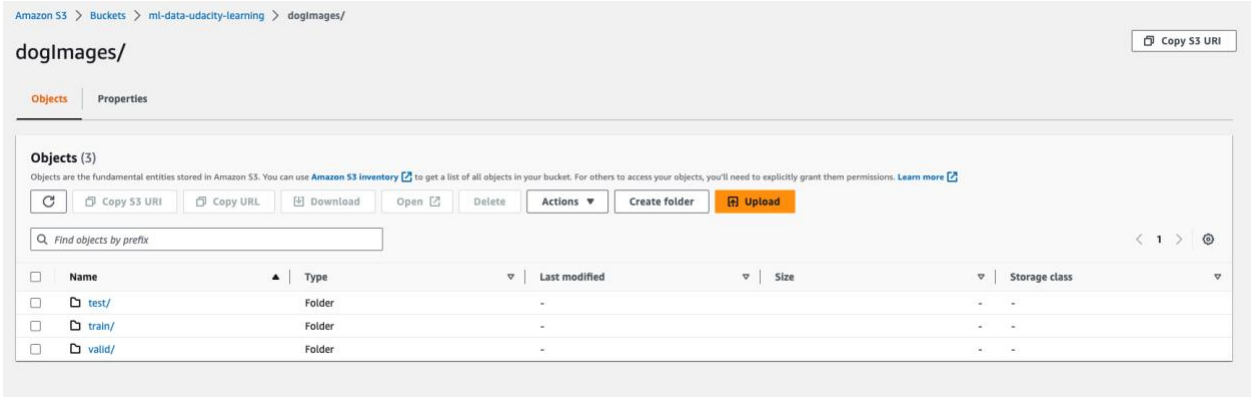

Single Instance Training Job

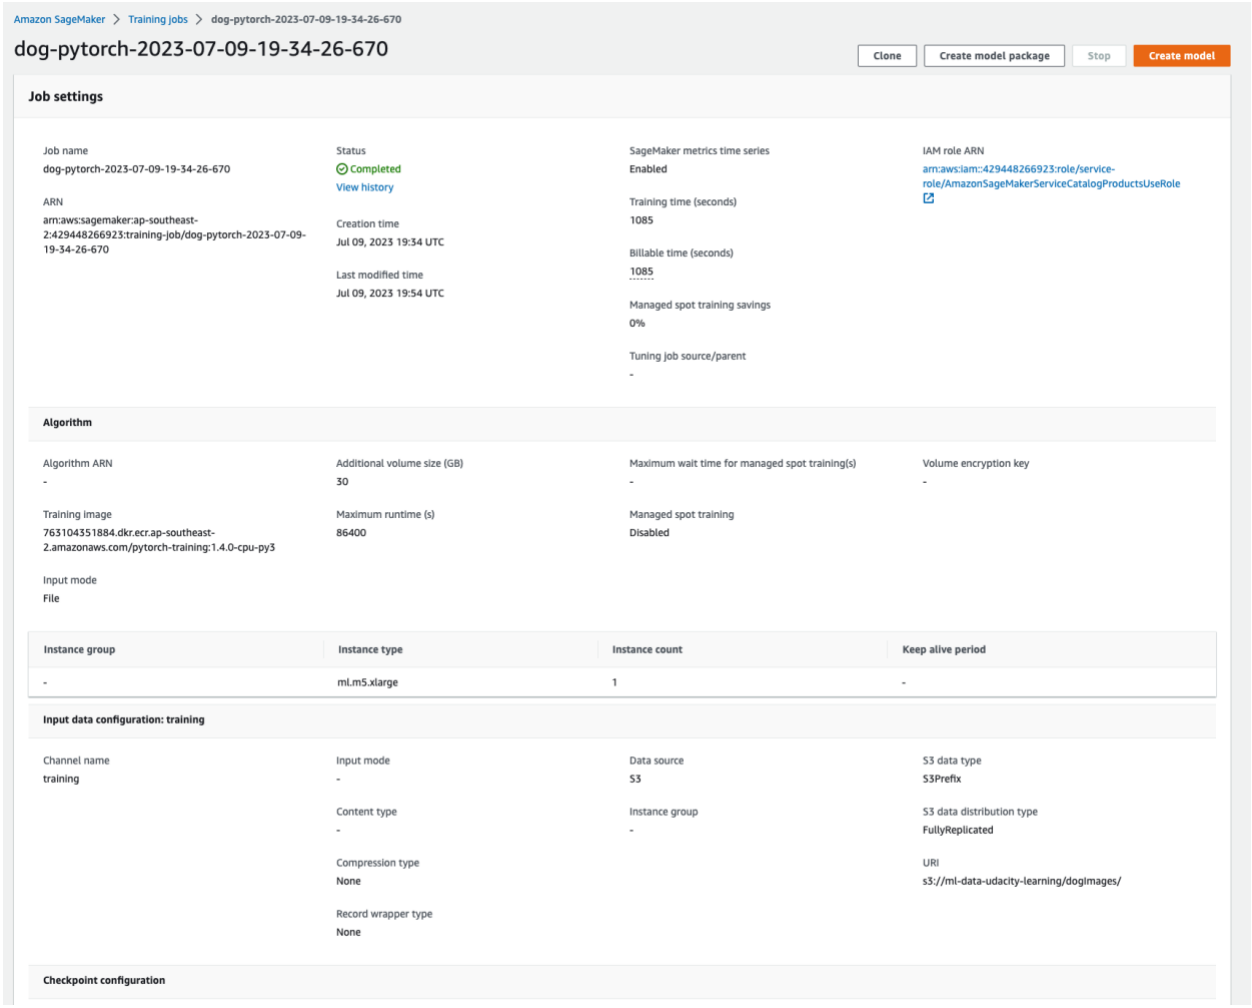

Both Training Jobs took fairly equal amount of time and I am very surprised.

Multi-Instance Training Job:

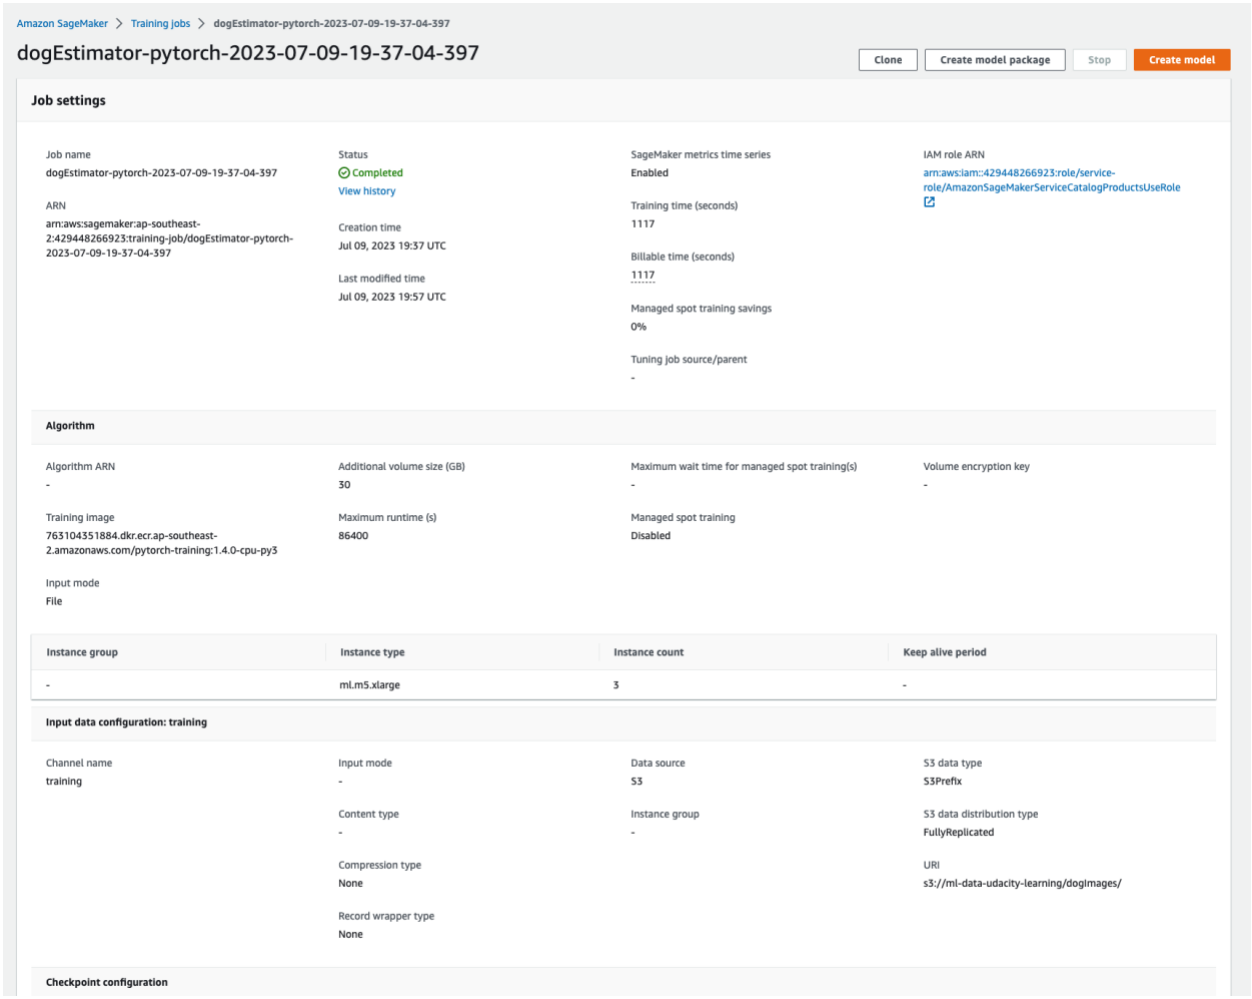

Deployment Endpoints

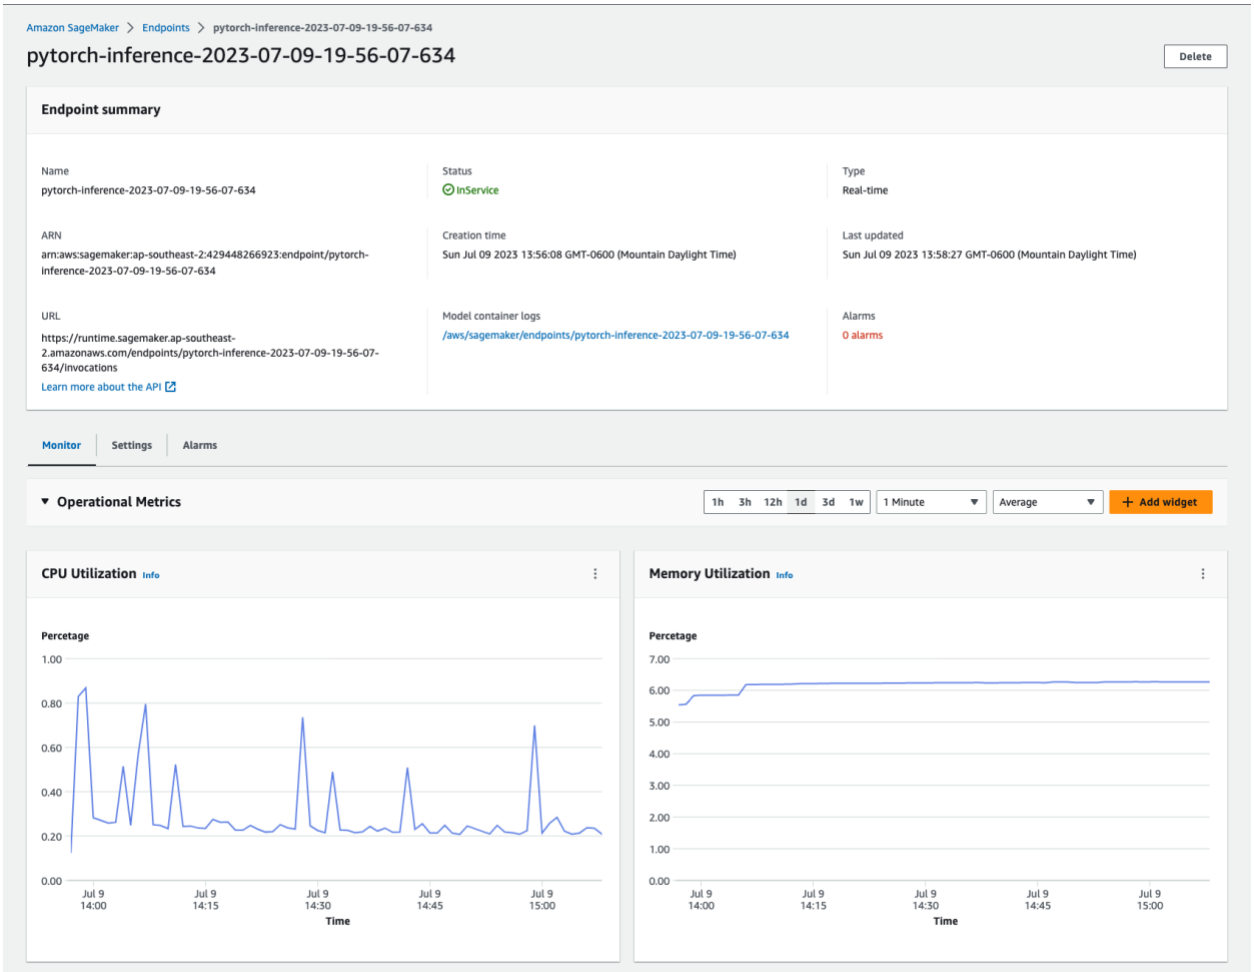
Amazon SageMaker > Endpoints > pytorch-inference-2023-07-09-19-56-07-634
pytorch-inference-2023-07-09-19-56-07-634
Delete
Endpoint summary
Name
pytorch-inference-2023-07-09-19-56-07-634
Status
InService
Type
Real-time
ARN
arn:aws:sagemaker:ap-southeast-2:429448266923:endpoint/pytorch-inference-2023-07-09-19-56-07-634
Creation time
Sun Jul 09 2023 13:56:08 GMT-0600 (Mountain Daylight Time)
Last updated
Sun Jul 09 2023 13:58:27 GMT-0600 (Mountain Daylight Time)
URL
https://runtime.sagemaker.ap-southeast-2.amazonaws.com/endpoints/pytorch-inference-2023-07-09-19-56-07-634/invocations
Learn more about the API
Model container logs
/aws/sagemaker/endpoints/pytorch-inference-2023-07-09-19-56-07-634
Alarms
0 alarms
Monitor
Settings
Alarms
Operational Metrics
1h 3h 12h 1d 3d 1w
1 Minute
Average
Add widget
CPU Utilization Info
Percetage
1.00
0.80
0.60
0.40
0.20
0.00
Jul 9 14:00
Jul 9 14:15
Jul 9 14:30
Jul 9 14:45
Jul 9 15:00
Time
Memory Utilization Info
Percetage
7.00
6.00
5.00
4.00
3.00
2.00
1.00
0.00
Jul 9 14:00
Jul 9 14:15
Jul 9 14:30
Jul 9 14:45
Jul 9 15:00
Time

```
-----!
```

```
[43]: import requests
      # request_dict={ "url": "https://cdn1-www.cattime.com/assets/uploads/2011/12/file_2744_british-shorthair-460x290-460x290.jpg" }
      request_dict={ "url": "https://s3.amazonaws.com/cdn-origin-etr.akc.org/wp-content/uploads/2017/11/20113314/Carolina-Dog-standing-outdoors.jpg" }

      img_bytes = requests.get(request_dict['url']).content
      type(img_bytes)
```

```
[43]: bytes
```

```
[44]: from PIL import Image
      import io
      Image.open(io.BytesIO(img_bytes))
```

```
[44]:
```

```
[31]: response=predictor.predict(img_bytes, initial_args={"ContentType": "image/jpeg"})
```

```
[32]: import json
      response2=predictor.predict(json.dumps(request_dict), initial_args={"ContentType": "application/json"})
```

```
[33]: type(response2[0][0])
```

```
[33]: float
```

```
[34]: response2[0]
```

```
[34]: [0.23838894069194794,
       0.09889235347509384,
       -0.132255420088768,
       0.34696149826049805,
       0.5369798541069031,
       0.25741326808929443,
       -0.1189652681350708,
       -0.009285075590014458,
       -0.35904017090797424,
```

2. EC2 Training

I employed the Deep Learning AMI, specifically GPU PyTorch 2.0.1 (Amazon Linux 2) 20230627, ami-051619310404cab17 (64-bit (x86)), alongside the g4dn.xlarge instance. This combination strikes an effective balance between financial considerations and performance capabilities.

To begin with, g4dn.xlarge instances emerge as the most economical option for this project, primarily due to the fact that the g4dn series generally represents the least expensive instances compatible with the Deep Learning AMI.

According to the official documentation, g4dn.xlarge instances are capable of maintaining superior CPU performance for the duration required by a task, a characteristic that adds to their suitability for this project.

Lastly, without incurring any additional expenses, g4dn instances provide adequate performance for the majority of applications of a general nature.

This instance type is supported by several EC2 instances, namely G3, P3, P3dn, P4d, P4de, G5, and G4dn. For further details, you can refer to the release notes available at: https://docs.aws.amazon.com/dlami/latest/devguide/appendix-ami-release-notes.html.

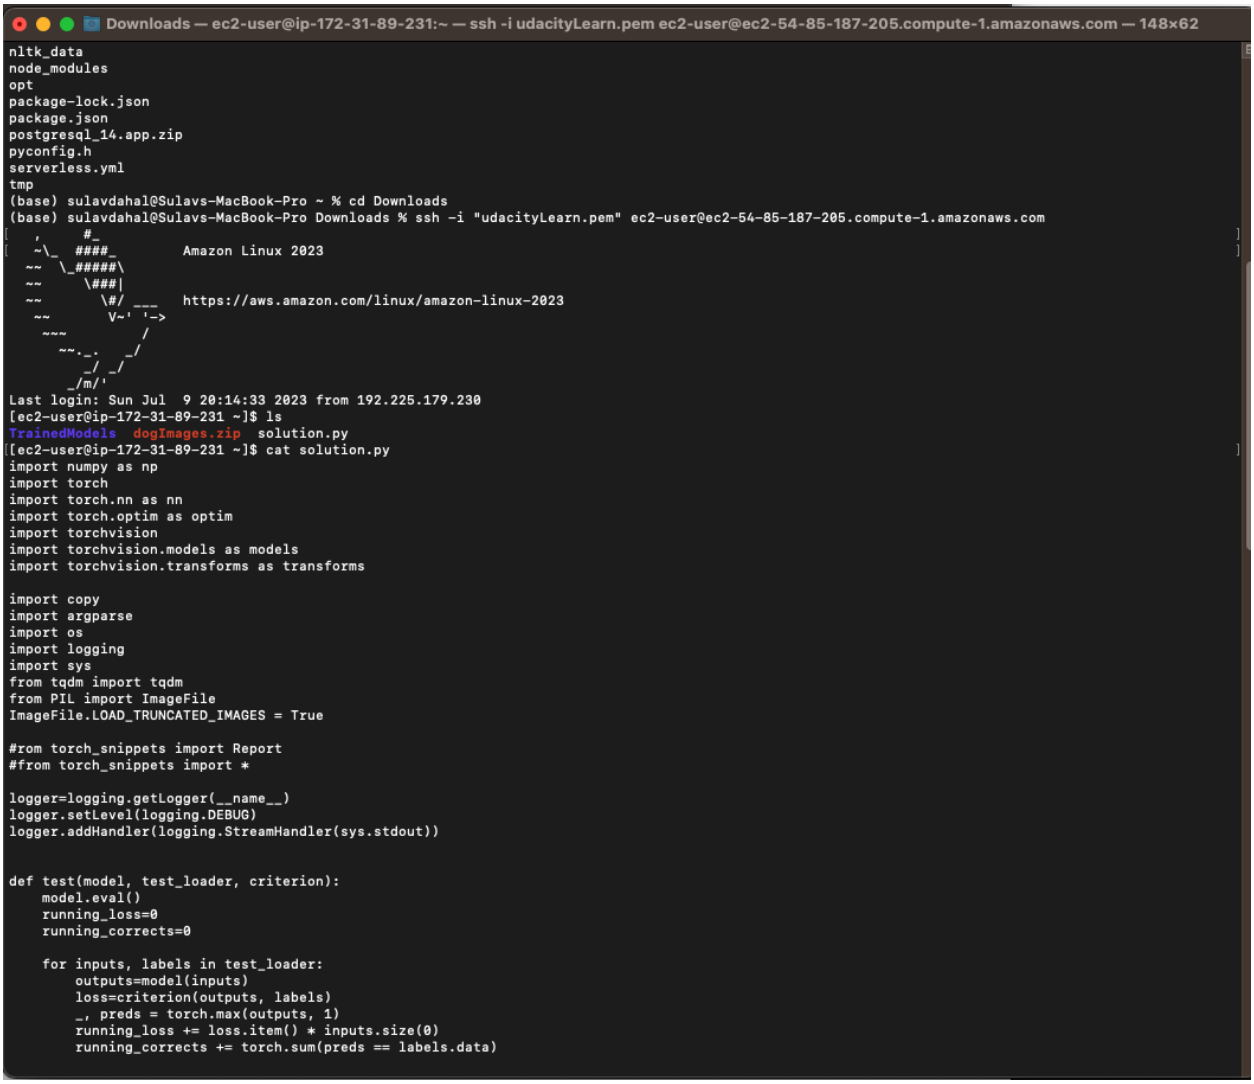

3. Lambda Function

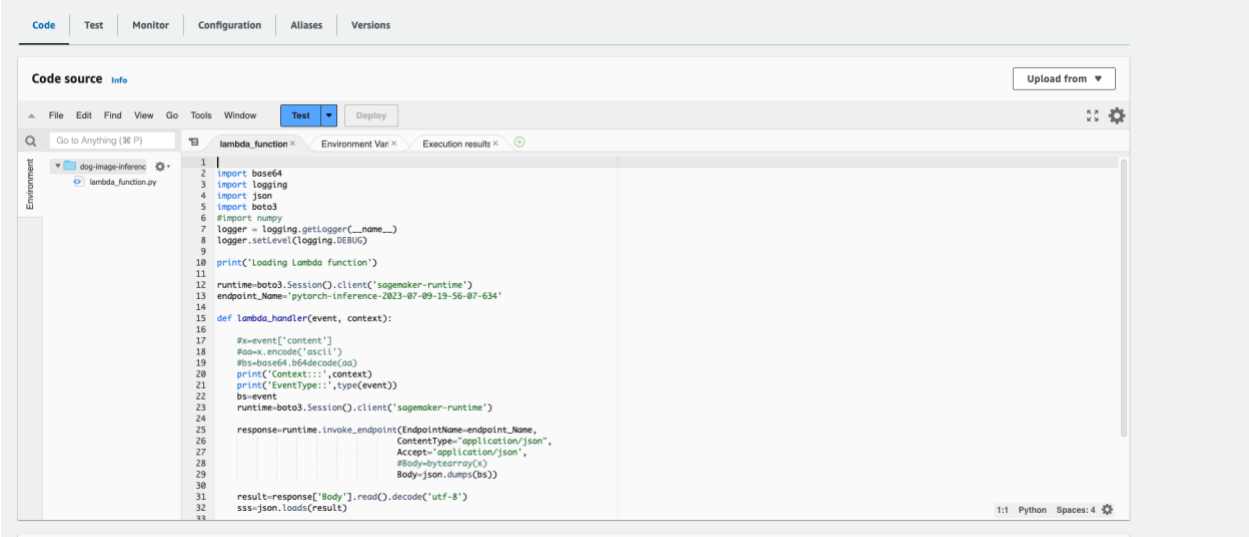

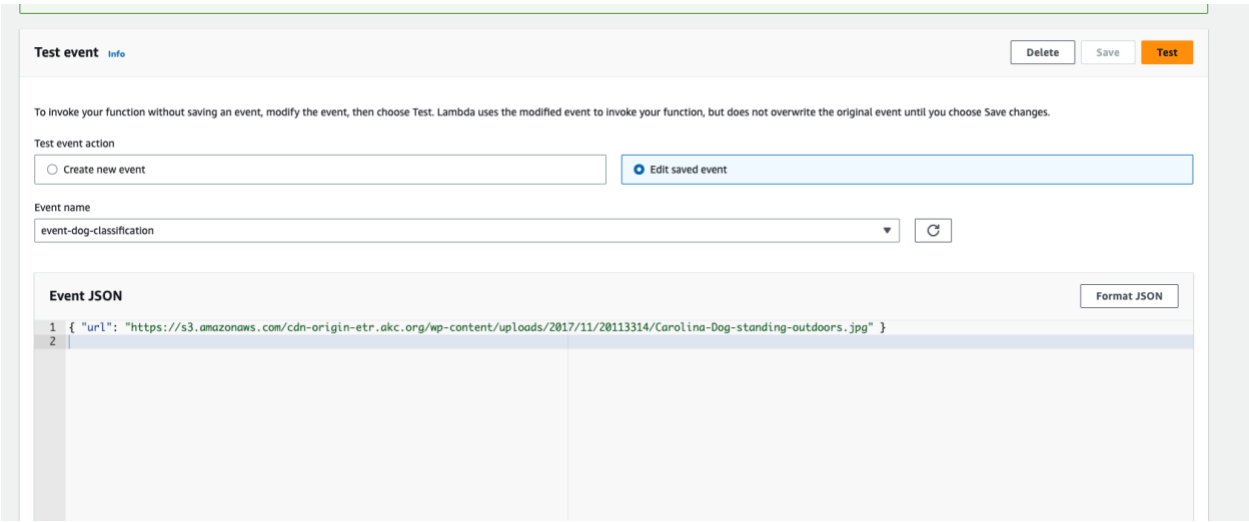

A SageMaker Policy has been attached to make the lambda code functional.

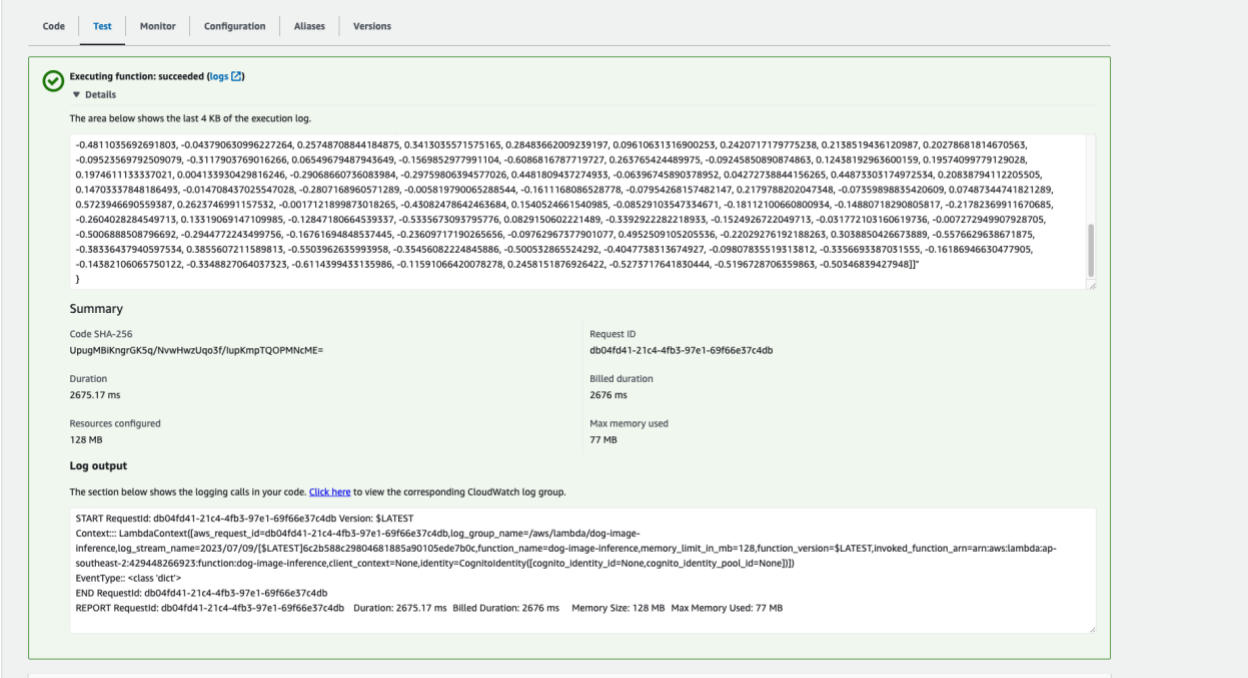

Concurrency

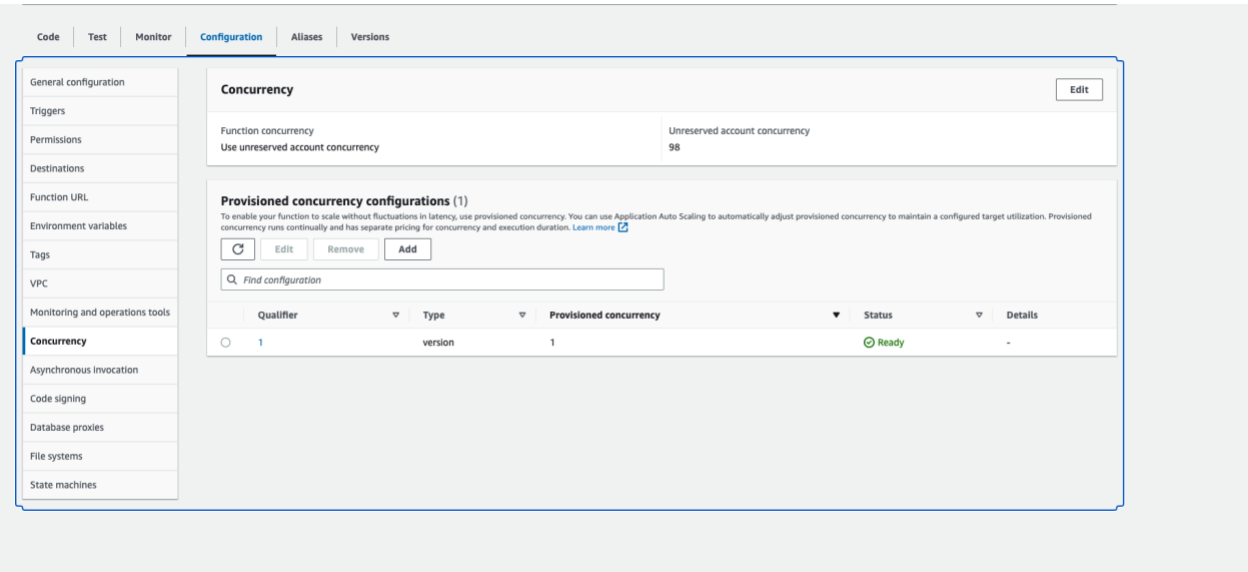

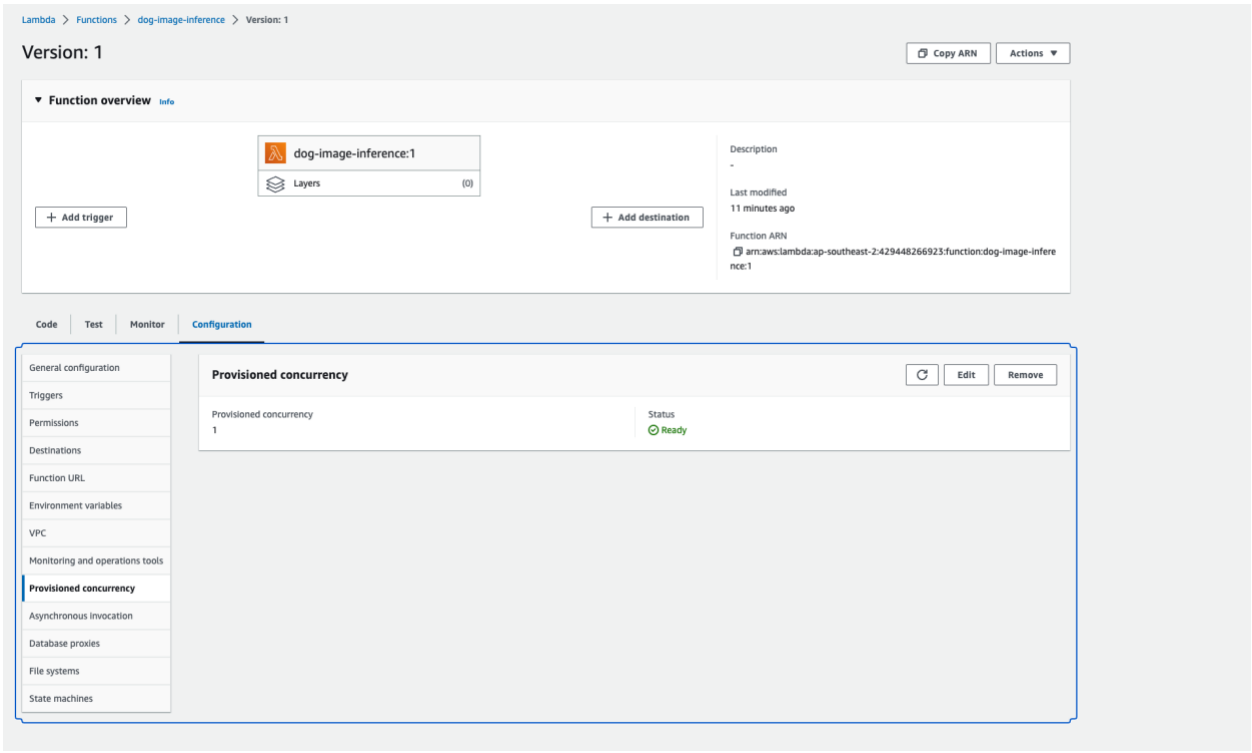
Lambda > Functions > dog-image-inference > Version: 1
Version: 1
Copy ARN
Actions
Function overview Info
dog-image-inference:1
Layers (0)
Add trigger
Add destination
Description
-
Last modified
11 minutes ago
Function ARN
arn:aws:lambda:ap-southeast-2:429448266923:function:dog-image-inference:1
Code
Test
Monitor
Configuration
General configuration
Triggers
Permissions
Destinations
Function URL
Environment variables
VPC
Monitoring and operations tools
Provisioned concurrency
Asynchronous invocation
Database proxies
File systems
State machines
Provisioned concurrency
Edit
Remove
Provisioned concurrency
1
Status
Ready